


Tuned Lens - Llama-3.2-3B on Mess3 HMM ($\alpha=0.85$, $x=0.05$)

KL Divergence: Logit Lens vs Tuned Lens



Belief-State R^2 vs Tuned-Lens KL per Layer

