



Tuned Lens — wing {'x': 0.98, 'y': 0.4}

KL Divergence by Layer



Belief-State R^2 vs Tuned-Lens KL per Layer



KL(model || lens): Model-target vs HMM-target

