



Tuned Lens — strata {'a': 0.9, 't0': 0.38, 't1': 0.54}

KL Divergence by Layer



Belief-State R^2 vs Tuned-Lens KL per Layer



KL(model || lens): Model-target vs HMM-target

