



KL Divergence by Layer

Belief-State R^2 vs Tuned-Lens KL per Layer

KL(model || lens): Model-target vs HMM-target

