



Tuned Lens — arch {'a': 0.91}

KL Divergence by Layer



Belief-State R^2 vs Tuned-Lens KL per Layer



KL(model || lens): Model-target vs HMM-target

