



Tuned Lens — mess3 {'a': 0.005, 'x': 0.01}

KL Divergence by Layer



Belief-State R^2 vs Tuned-Lens KL per Layer



KL(model || lens): Model-target vs HMM-target

