



Tuned Lens — mess3 {'x': 0.05, 'a': 0.85}

KL Divergence by Layer



Belief-State R^2 vs Tuned-Lens KL per Layer



KL(model || lens): Model-target vs HMM-target

